

Introduction

Mobile phones come in all sorts of prices, features, specifications and all. Price estimation and prediction is an important part of consumer strategy. Deciding on the correct price of a product is very important for the market success of a product. A new product that has to be launched, must have the correct price so that consumers find it appropriate to buy the product.

Problem Statement

To predict the price range of a mobile phone. The data contains information regarding mobile phone features, specifications etc and their price range. The various features and information can be used to predict the price range of a mobile phone.

```
In [2]: # import libraries
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings('ignore')
```

Read data

```
In [3]: data=pd.read_csv(r"C:\Users\samee\Downloads\ttrain_data.csv")
data.head()
```

```
Out[3]:
```

	id	battery_power	bluetooth	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	px_height	px_width	ram	sc_h	sc
0	1	807	1	0.5	1	0	0	37	0.2	127	...	245	829	2319	5	
1	2	1065	1	0.5	0	0	1	14	0.7	89	...	188	928	3078	10	
2	3	1171	1	1.7	1	2	0	19	0.3	167	...	248	755	263	6	
3	4	609	1	3.0	0	15	1	44	0.3	117	...	58	1253	2581	15	
4	5	1193	1	2.3	0	7	0	20	1.0	158	...	1442	1904	1958	7	

5 rows × 22 columns

```
In [4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1400 entries, 0 to 1399
Data columns (total 22 columns):
#   Column              Non-Null Count  Dtype
---  -
0   id                  1400 non-null   int64
1   battery_power       1400 non-null   int64
2   bluetooth           1400 non-null   int64
3   clock_speed         1400 non-null   float64
4   dual_sim            1400 non-null   int64
5   fc                  1400 non-null   int64
6   four_g              1400 non-null   int64
7   int_memory          1400 non-null   int64
8   m_dep               1400 non-null   float64
9   mobile_wt           1400 non-null   int64
10  n_cores              1400 non-null   int64
11  pc                   1400 non-null   int64
12  px_height            1400 non-null   int64
13  px_width             1400 non-null   int64
14  ram                  1400 non-null   int64
15  sc_h                 1400 non-null   int64
16  sc_w                 1400 non-null   int64
17  talk_time            1400 non-null   int64
18  three_g              1400 non-null   int64
19  touch_screen        1400 non-null   int64
20  wifi                 1400 non-null   int64
21  price_range          1400 non-null   int64
dtypes: float64(2), int64(20)
memory usage: 240.8 KB
```

```
In [5]: data.shape
```

```
Out[5]: (1400, 22)
```

```
In [6]: data.price_range.value_counts()
```

```
Out[6]: 1    350
        2    350
        0    350
        3    350
        Name: price_range, dtype: int64
```

Cleaning Part

```
In [7]: l=data["id"]
        l
```

```
Out[7]: 0      1
        1      2
        2      3
        3      4
        4      5

        ...
        1395    1396
        1396    1397
        1397    1398
        1398    1399
        1399    1400
        Name: id, Length: 1400, dtype: int64
```

```
In [8]: mob_data=data.drop(['id'], axis=1)
        mob_data.head()
```

```
Out[8]:
```

	battery_power	bluetooth	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc
0	807	1	0.5	1	0	0	37	0.2	127	4	...	245	829	2319	
1	1065	1	0.5	0	0	1	14	0.7	89	2	...	188	928	3078	
2	1171	1	1.7	1	2	0	19	0.3	167	7	...	248	755	263	
3	609	1	3.0	0	15	1	44	0.3	117	1	...	58	1253	2581	
4	1193	1	2.3	0	7	0	20	1.0	158	7	...	1442	1904	1958	

5 rows × 21 columns

Missing Value Treatment

Missing values are usually represented in the form of Nan or null or None in the dataset.

```
In [9]: mob_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1400 entries, 0 to 1399
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype  
---  -
0   battery_power    1400 non-null   int64  
1   bluetooth        1400 non-null   int64  
2   clock_speed      1400 non-null   float64 
3   dual_sim         1400 non-null   int64  
4   fc               1400 non-null   int64  
5   four_g           1400 non-null   int64  
6   int_memory       1400 non-null   int64  
7   m_dep            1400 non-null   float64 
8   mobile_wt        1400 non-null   int64  
9   n_cores          1400 non-null   int64  
10  pc               1400 non-null   int64  
11  px_height        1400 non-null   int64  
12  px_width         1400 non-null   int64  
13  ram              1400 non-null   int64  
14  sc_h             1400 non-null   int64  
15  sc_w             1400 non-null   int64  
16  talk_time        1400 non-null   int64  
17  three_g          1400 non-null   int64  
18  touch_screen     1400 non-null   int64  
19  wifi             1400 non-null   int64  
20  price_range      1400 non-null   int64  
dtypes: float64(2), int64(19)
memory usage: 229.8 KB
```

```
In [10]: mob_data.isnull().sum()
```

```
Out[10]: battery_power      0
         bluetooth         0
         clock_speed       0
         dual_sim          0
         fc                0
         four_g            0
         int_memory        0
         m_dep             0
         mobile_wt         0
         n_cores           0
         pc                0
         px_height         0
         px_width          0
         ram               0
         sc_h              0
         sc_w              0
         talk_time         0
         three_g           0
         touch_screen      0
         wifi              0
         price_range       0
         dtype: int64
```

```
In [11]: #Remove the data points with missing data
         mob_data_f = mob_data[mob_data['sc_w'] != 0]
         mob_data_f.shape
```

```
Out[11]: (1276, 21)
```

Data types Conversion

```
In [12]: mob_data_f.dtypes
```

```
Out[12]: battery_power      int64
         bluetooth         int64
         clock_speed       float64
         dual_sim          int64
         fc                int64
         four_g            int64
         int_memory        int64
         m_dep             float64
         mobile_wt         int64
         n_cores           int64
         pc                int64
         px_height         int64
         px_width          int64
         ram               int64
         sc_h              int64
         sc_w              int64
         talk_time         int64
         three_g           int64
         touch_screen      int64
         wifi              int64
         price_range       int64
         dtype: object
```

Label Encoder

Label Encoding refers to converting the labels into a numeric form so as to convert them into the machine-readable form.

```
In [13]: from sklearn.preprocessing import LabelEncoder
         le=LabelEncoder()
```

```
In [14]: mob_data_f.battery_power=le.fit_transform(mob_data_f.battery_power)
```

```
In [15]: mob_data_f.bluetooth=le.fit_transform(mob_data_f.bluetooth)
```

```
In [16]: mob_data_f.battery_power.value_counts()
```

```
Out[16]: 794      5
         77       4
         287      4
         627      4
         203      4
         ..
         316      1
         591      1
         119      1
         596      1
         333      1
         Name: battery_power, Length: 874, dtype: int64
```

```
In [17]: mob_data_f.bluetooth.value_counts()
```

```
In [17]: mob_data_f.bluetooth_count()
Out[17]: 1    647
         0    629
         Name: bluetooth, dtype: int64
```

```
In [18]: mob_data_f.dtypes
```

```
Out[18]: battery_power    int64
         bluetooth        int64
         clock_speed      float64
         dual_sim         int64
         fc               int64
         four_g           int64
         int_memory       int64
         m_dep            float64
         mobile_wt        int64
         n_cores          int64
         pc              int64
         px_height        int64
         px_width         int64
         ram              int64
         sc_h             int64
         sc_w             int64
         talk_time        int64
         three_g          int64
         touch_screen     int64
         wifi             int64
         price_range      int64
         dtype: object
```

Exploratory Data Analysis(EDA)

```
In [19]: mob_data_f.skew()
```

```
Out[19]: battery_power    0.023699
         bluetooth       -0.028249
         clock_speed      0.158978
         dual_sim         0.000000
         fc              1.061322
         four_g          -0.081669
         int_memory       0.066551
         m_dep            0.084648
         mobile_wt       -0.030716
         n_cores          -0.029976
         pc              0.045808
         px_height        0.654526
         px_width         0.031901
         ram             -0.000004
         sc_h            -0.128919
         sc_w            0.662084
         talk_time        0.006666
         three_g         -1.205159
         touch_screen     -0.037669
         wifi            -0.069088
         price_range      -0.016020
         dtype: float64
```

Correlation

```
In [20]: x=mob_data_f.corr()
         x
```


Out[20]:


	battery_power	bluetooth	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_h
battery_power	1.000000	-0.003565	0.005420	-0.046991	0.021736	0.028544	0.006901	0.044438	0.021853	-0.017810	...	-0.01
bluetooth	-0.003565	1.000000	0.043958	-0.001568	-0.014205	-0.011557	0.022142	-0.004980	-0.026344	0.051186	...	-0.01
clock_speed	0.005420	0.043958	1.000000	0.004006	-0.022518	-0.008297	-0.000558	0.005163	0.003277	-0.007766	...	-0.00
dual_sim	-0.046991	-0.001568	0.004006	1.000000	-0.012710	-0.009412	-0.002372	0.002153	0.000553	-0.002415	...	-0.00
fc	0.021736	-0.014205	-0.022518	-0.012710	1.000000	-0.018963	-0.040908	-0.016456	0.011467	-0.028517	...	0.01
four_g	0.028544	-0.011557	-0.008297	-0.009412	-0.018963	1.000000	0.025071	-0.055330	0.021100	-0.029581	...	-0.00
int_memory	0.006901	0.022142	-0.000558	-0.002372	-0.040908	0.025071	1.000000	-0.017206	-0.060900	-0.020671	...	-0.03
m_dep	0.044438	-0.004980	0.005163	0.002153	-0.016456	-0.055330	-0.017206	1.000000	0.032231	-0.004648	...	0.04
mobile_wt	0.021853	-0.026344	0.003277	0.000553	0.011467	0.021100	-0.060900	0.032231	1.000000	-0.019822	...	0.02
n_cores	-0.017810	0.051186	-0.007766	-0.002415	-0.028517	-0.029581	-0.020671	-0.004648	-0.019822	1.000000	...	0.02
pc	0.036868	-0.008959	-0.010164	-0.025674	0.640535	0.005825	-0.050473	0.019213	0.018134	-0.020636	...	-0.00
px_height	-0.013950	-0.010005	-0.000148	-0.001552	0.014964	-0.007069	-0.030512	0.047354	0.024427	0.020605	...	1.00
px_width	-0.007997	-0.028540	0.002970	0.036019	0.005239	0.005759	-0.020931	0.041478	0.006862	0.053240	...	0.50
ram	0.001132	0.019699	-0.018746	0.062861	-0.006512	-0.024694	0.025228	-0.012488	-0.009125	-0.017298	...	-0.01
sc_h	-0.028019	-0.009435	-0.042018	-0.005545	-0.007226	0.038833	0.023074	-0.025123	-0.020312	-0.025482	...	0.07
sc_w	-0.018282	-0.016202	-0.009298	-0.011640	0.015149	0.024770	-0.009918	-0.007579	-0.035125	0.000893	...	0.03
talk_time	0.058074	0.018587	-0.011112	-0.014444	-0.011048	-0.065537	-0.008219	0.016651	-0.023781	-0.002123	...	-0.00
three_g	0.006225	-0.045083	-0.026799	-0.034759	0.003212	0.588809	0.006206	-0.038490	0.014318	-0.006161	...	-0.03
touch_screen	-0.003661	0.004438	0.052383	0.009406	-0.022933	0.011785	-0.002606	-0.002873	-0.016183	0.045139	...	0.06
wifi	0.005654	-0.008329	-0.011893	0.025093	0.010131	-0.035939	0.006445	-0.035872	-0.012720	-0.009307	...	0.04
price_range	0.195348	0.006854	-0.021327	0.037107	-0.002291	-0.012340	0.039866	0.010475	-0.035872	-0.005655	...	0.14

21 rows × 21 columns


In [21]:

```
plt.figure(figsize=(20,30))
heatmap=sns.heatmap(x,cmap="coolwarm",annot=True)
plt.title("Heatmap for correlation")
plt.show()
```






```
In [22]: plt.figure(figsize=(20,9))
sns.boxplot(data=mob_data_f)
plt.xticks(rotation=70)
plt.show()
```



```
In [23]: sns.set()
price_plot=mob_data_f['price_range'].value_counts().plot(kind='bar')
plt.xlabel('price_range')
plt.ylabel('Count')
plt.show()
```




So, there are mobile phones in 4 price ranges. The number of elements is almost similar.

Data Distribution


First, we see how the battery_power mAh is spread.

```
In [24]: sns.set(rc={'figure.figsize': (5,5)})  
ax=sns.displot(data=mob_data_f["battery_power"])  
plt.show()
```




Now, we see the count of how many devices have Bluetooth and how many don't.


```
In [25]: sns.distplot(mob_data_f.battery_power)  
plt.show()
```




```
In [26]: sns.boxplot(data=mob_data_f,y='battery_power')  
plt.show()
```



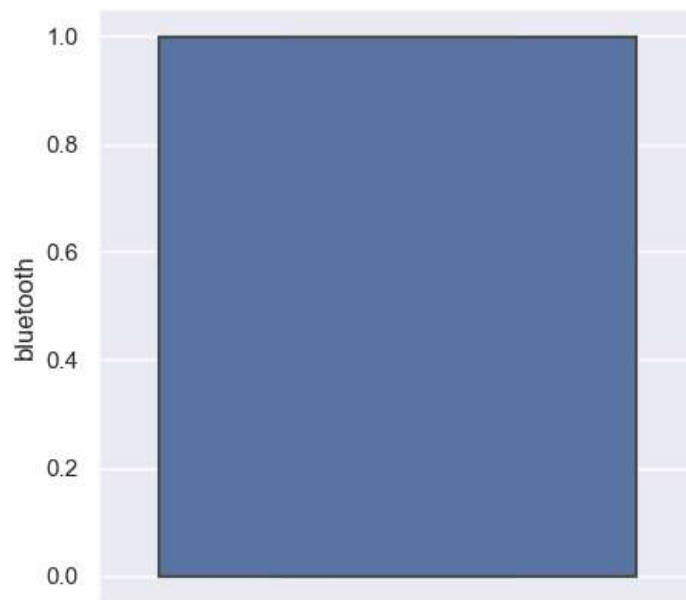
```
In [27]: sns.set(rc={'figure.figsize':(5,5)})  
ax=sns.displot(data=mob_data_f["bluetooth"])  
plt.show()
```

```
In [28]: sns.distplot(mob_data_f.bluetooth)
plt.show()
```




```
In [29]: sns.boxplot(data=mob_data_f,y='bluetooth')
plt.show()
```




Next, we analyse the mobile depth (in cm).

```
In [30]: sns.set(rc={'figure.figsize':(5,5)})  
ax=sns.displot(data=mob_data_f["m_dep"])  
plt.show()
```




-----> A few mobiles are very thin and a few ones are almost a cm thick.

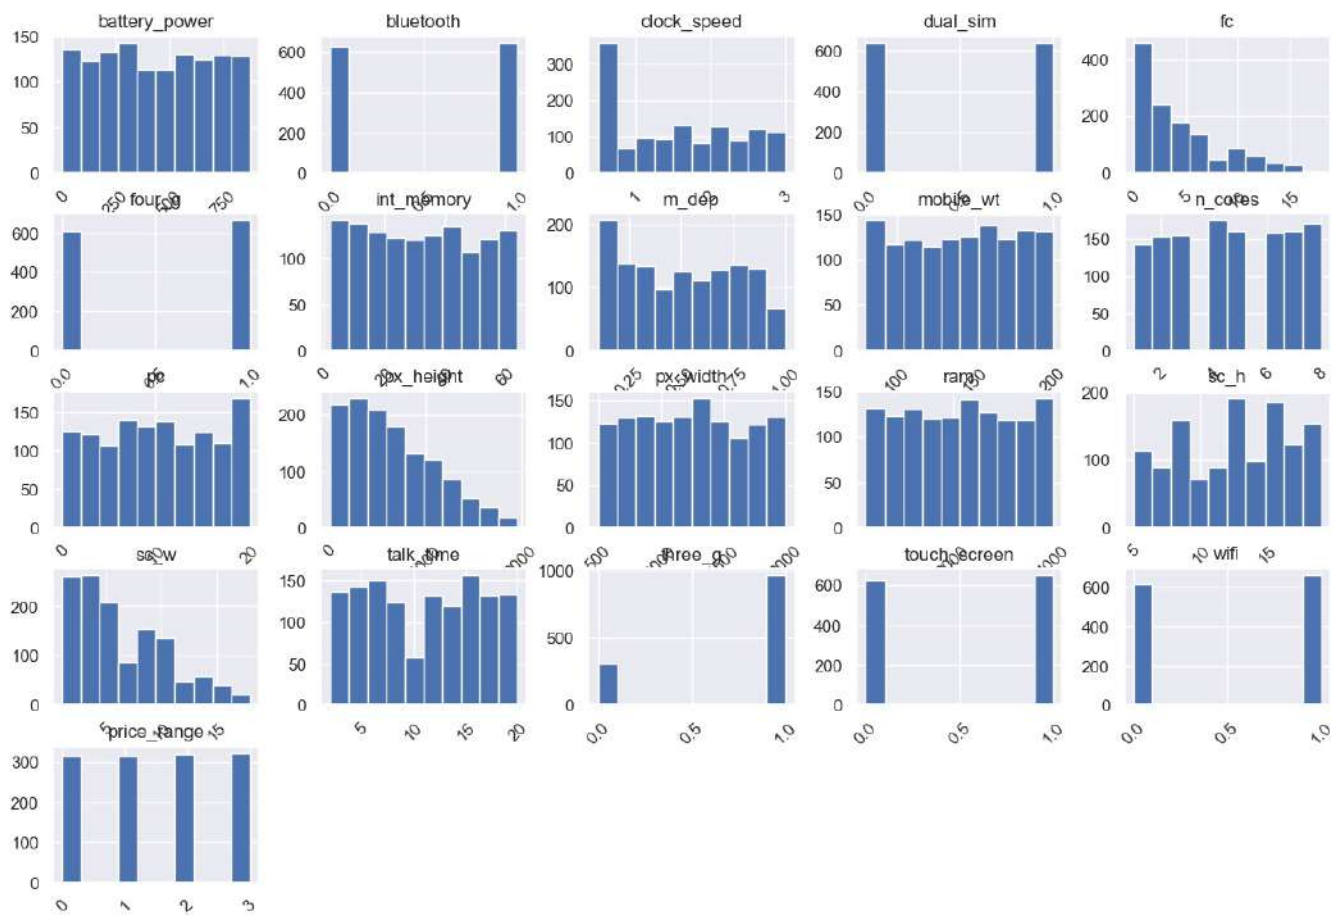
```
In [31]: sns.distplot(mob_data_f.m_dep)  
plt.show()
```



```
In [32]: sns.boxplot(data=mob_data_f,y='m_dep')  
plt.show()
```



```
In [33]: mob_data_f.hist(figsize=(15,10),xrot=45)  
plt.xticks(rotation=70)  
plt.show()
```



Creating Base Model

In a similar way, the data distribution can be analysed for all the data features. Implementing that will be very simple. Let us see if there are any missing values or missing data.

```
In [34]: mob_data_f.shape
```

```
Out[34]: (1276, 21)
```

```
In [35]: x=mob_data_f.drop(['price_range'], axis=1)
y=mob_data_f['price_range']
# missing values
x.isna().any()
```

```
Out[35]: battery_power    False
bluetooth              False
clock_speed            False
dual_sim               False
fc                     False
four_g                 False
int_memory             False
m_dep                  False
mobile_wt              False
n_cores                False
pc                     False
px_height              False
px_width               False
ram                    False
sc_h                   False
sc_w                   False
talk_time              False
three_g                False
touch_screen           False
wifi                   False
dtype: bool
```

Split the data

```
In [66]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, random_state=7)
```

Now, we define a function for creating confusion matrix.

```
In [68]: #Function for confusion matrix
```

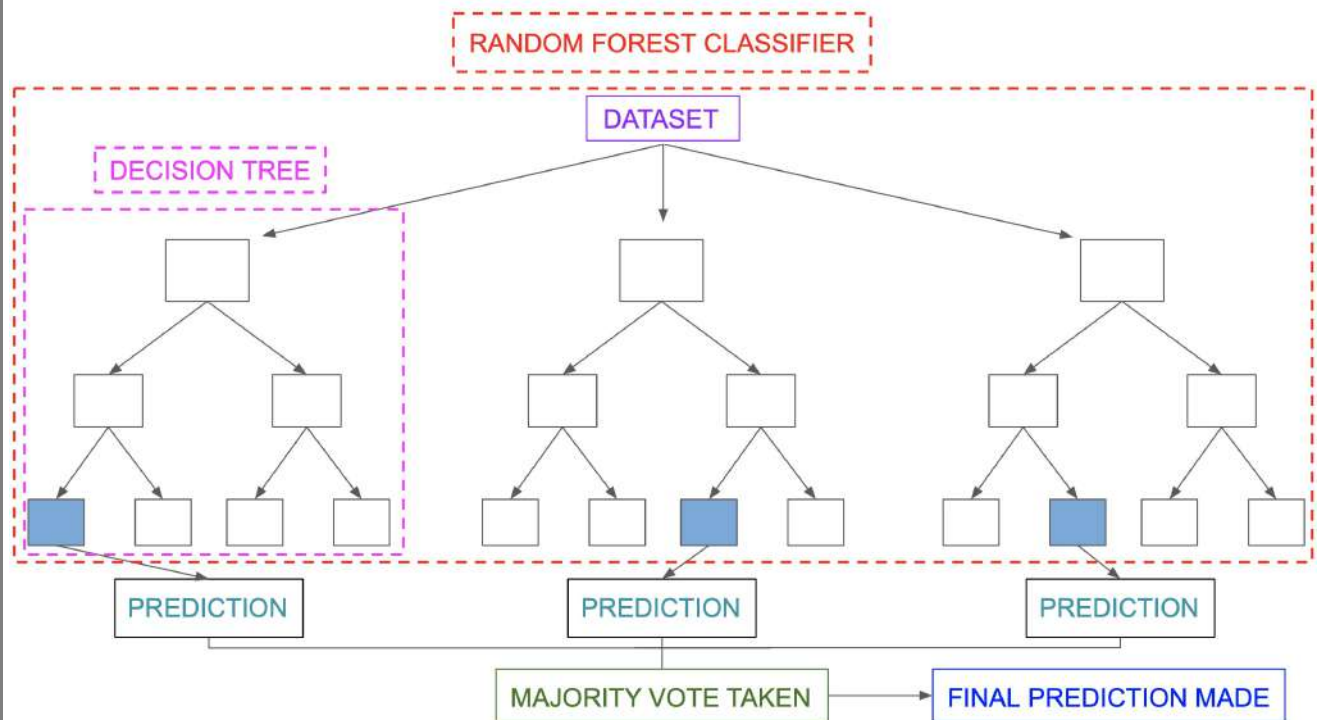
```

from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
def my_confusion_matrix(y_test, y_pred, plt_title):
    cm=confusion_matrix(y_test, y_pred)
    print(classification_report(y_test, y_pred))
    sns.heatmap(cm, annot=True, fmt='g', cbar=False, cmap='BuPu')
    plt.xlabel('Predicted Values')
    plt.ylabel('Actual Values')
    plt.title(plt_title)
    plt.show()
    return cm

```

1) Random Forest Classifier

A random forest is a supervised machine learning method built from decision tree techniques. This algorithm is used to anticipate behaviour and results in a variety of sectors, including banking and e-commerce. A random forest is a machine learning approach for solving regression and classification issues. It makes use of ensemble learning, which is a technique that combines multiple classifiers to solve complicated problems. The Random Forest Algorithm based on Decision Tree Algorithm



```

In [72]: #Building the model
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier(bootstrap=True,
                           max_depth= 7,
                           max_features= 15,
                           min_samples_leaf= 3,
                           min_samples_split= 10,
                           n_estimators= 200,
                           random_state=7)
rfc.fit(x_train, y_train)

```

```

Out[72]: ▼      RandomForestClassifier
RandomForestClassifier(max_depth=7, max_features=15, min_samples_leaf=3,
                       min_samples_split=10, n_estimators=200, random_state=7)

```

```

In [73]: y_pred_rfc=rfc.predict(x_test)

```

```

In [74]: print('Random Forest Classifier Accuracy Score: ',accuracy_score(y_test,y_pred_rfc))
cm_rfc=my_confusion_matrix(y_test, y_pred_rfc, 'Random Forest Confusion Matrix')

```


```

Random Forest Classifier Accuracy Score: 0.90234375
precision    recall  f1-score   support

0           0.94     0.97     0.95         64
1           0.90     0.87     0.89         71
2           0.82     0.86     0.84         59
3           0.95     0.90     0.93         62

accuracy          0.90
macro avg          0.90
weighted avg       0.90

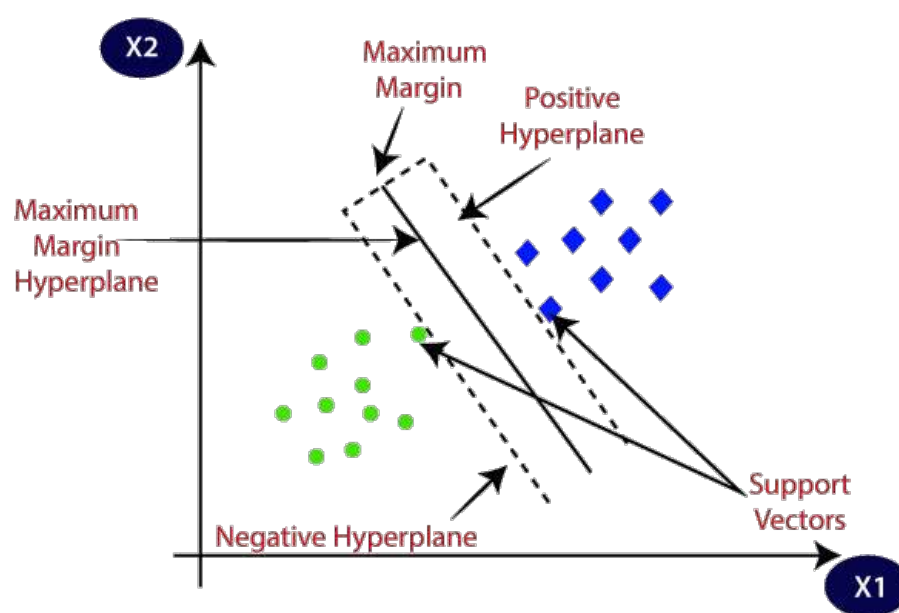
```



2) SVM Classifier

Support Vector Machine, or SVM, is a prominent Supervised Learning technique that is used for both classification and regression issues. However, it is mostly utilised in Machine Learning for Classification purposes.

The SVM algorithm's purpose is to find the optimum line or decision boundary for categorising n-dimensional space so that we may simply place fresh data points in the proper category in the future. A hyperplane is the optimal choice boundary.




```
In [75]: from sklearn import svm
svm_clf = svm.SVC(decision_function_shape='ovo')
svm_clf.fit(x_train, y_train)
```

```
Out[75]: SVC
SVC(decision_function_shape='ovo')
```

```
In [76]: y_pred_svm=svm_clf.predict(x_test)
```

```
In [77]: print('SVM Classifier Accuracy Score: ',accuracy_score(y_test,y_pred_svm))
cm_svm=my_confusion_matrix(y_test, y_pred_svm, 'SVM Confusion Matrix')
```


SVM Classifier Accuracy Score: 0.80859375					
	precision	recall	f1-score	support	
0	0.92	0.88	0.90	64	
1	0.75	0.76	0.76	71	
2	0.68	0.73	0.70	59	
3	0.90	0.87	0.89	62	
accuracy			0.81	256	
macro avg	0.81	0.81	0.81	256	
weighted avg	0.81	0.81	0.81	256	



We can see that the SVM classifier is giving the best accuracy.

3) Naive Bayes

Conditional probability is the foundation of Bayes' theorem. The conditional probability aids us in assessing the likelihood of something occurring if something else has previously occurred. Gaussian Naive Bayes is a Naive Bayes variation that allows continuous data and follows the Gaussian normal distribution.



z-score distance of x
from **Class A**

z-score distance of x
from **Class B**

```
In [54]: from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(x_train, y_train)
```

```
Out[54]: GaussianNB
GaussianNB()
```


```
In [55]: y_pred_gnb=gnb.predict(x_test)
```

```
In [56]: print('Gaussian NB Classifier Accuracy Score: ',accuracy_score(y_test,y_pred_gnb))
cm_gnb=my_confusion_matrix(y_test, y_pred_gnb, 'Gaussian NB Confusion Matrix')
```

```
Gaussian NB Classifier Accuracy Score: 0.8046875
      precision    recall  f1-score   support

     0       0.90      0.86      0.88         64
     1       0.74      0.73      0.74         71
     2       0.69      0.73      0.71         59
     3       0.89      0.90      0.90         62


 accuracy          0.80         256
 macro avg         0.81         256
 weighted avg      0.81         256
```




We can see that the model is performing well.

4) Decision Tree Classifier

The decision tree classifier creates the classification model by building a decision tree. Each node in the tree specifies a test on an attribute, each branch descending from that node corresponds to one of the possible values for that attribute.





```
In [57]: from sklearn.tree import DecisionTreeClassifier
dtc = DecisionTreeClassifier()
dtc.fit(x_train,y_train)
```

```
Out[57]: DecisionTreeClassifier
DecisionTreeClassifier()
```

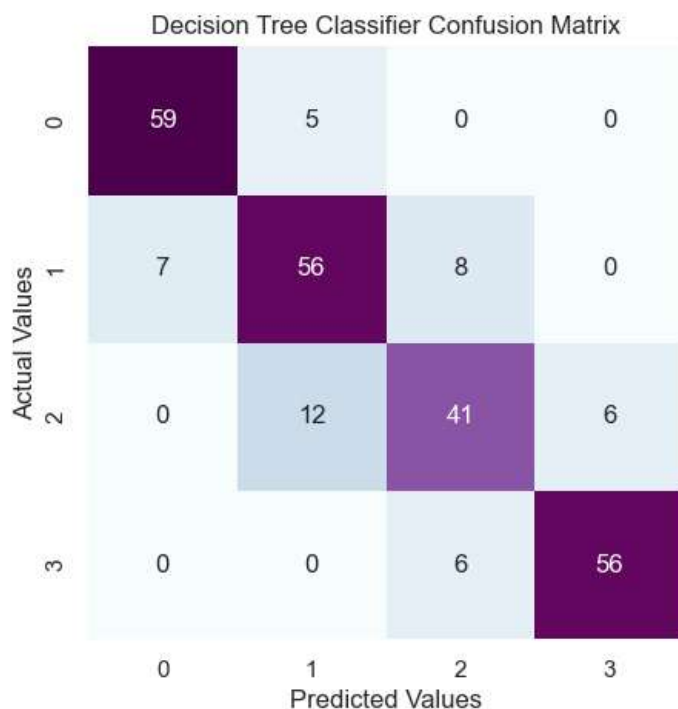
```
In [59]: y_pred_dtc = dtc.predict(x_test)
```

```
In [60]: print('Decision Tree Classifier Accuracy Score: ',accuracy_score(y_test,y_pred_dtc))
cm_dtc=my_confusion_matrix(y_test, y_pred_dtc, 'Decision Tree Classifier Confusion Matrix')
```

```
Decision Tree Classifier Accuracy Score: 0.828125
precision    recall  f1-score   support

     0       0.89     0.92     0.91         64
     1       0.77     0.79     0.78         71
     2       0.75     0.69     0.72         59
     3       0.90     0.90     0.90         62

 accuracy          0.83
 macro avg         0.83
weighted avg         0.83
```



Conclusion

In this problem, we looked at classification. Classifiers represent the intersection of advanced machine theory and practical application. These algorithms are more than just a sorting mechanism for organizing unlabeled data instances into distinct groupings. Classifiers include a unique set of dynamic rules that include an interpretation mechanism for dealing with ambiguous or unknown values, all of which are suited to the kind of inputs being analysed. Most classifiers also utilize probability estimates, which enable end-users to adjust data categorization using utility functions.

In this problem, we see that SVM Classifier is perform well than other models.

```
In [ ]:
```